PDF RAG Workspace: Public Demo Guide
Page 1: Finding relevant text

PDF RAG Workspace lets you read a PDF beside a chat and ask questions grounded in the document. The server extracts selectable text from each physical PDF page and chunks each page separately. MiniLM converts chunks into 384-dimensional embeddings. PostgreSQL with pgvector compares a question embedding with stored chunk embeddings using cosine similarity. Retrieval is restricted to the selected document owned by the authenticated account. The server sends the retrieved text to Groq as context and asks it to answer only from that context. Answers stream back to the browser using server-sent events over a fetch request. If the context does not provide an answer, the model is instructed to say it could not find the answer. Model answers can still be wrong; check the supporting passage.

Answers cite retrieved chunks using Chunk markers. Each source stores a chunk index, its physical PDF page number, a text preview and retrieval similarity. Clicking a citation opens the corresponding page in the PDF viewer and highlights the supporting text when a match is found. If matching fails, the viewer briefly highlights the page. The viewer uses PDF.js through react-pdf and lazy-renders pages. Citations refer to physical page positions, not printed page labels. Each document has a separate conversation. Chat history and drafts stay in the visitor’s browser; history is saved in localStorage, while drafts stay in memory. Conversations are not synced between devices. The interface supports Markdown answers, light and dark themes, a desktop split view and mobile chat/PDF tabs.

PDF RAG Workspace: Public Demo Guide
Page 3: Guest demo and full accounts

This is a public sample PDF created specifically for the guest demo. Guest access requires no email, password or registration. A guest session lasts 20 minutes and allows reading this sample and asking questions about it. Guests cannot upload PDFs, rename or delete this sample, or access other people’s documents. All guest questions share an allowance of 10 requests per minute on the current single backend process.
To upload your own PDF, sign in or create a free account. Full accounts have owner-scoped document access, uploads up to 10 MiB, rename and delete. Ingestion is synchronous, so the UI reports real upload progress followed by indeterminate processing. Scanned PDFs without selectable text need OCR, which is not implemented.
The frontend runs on Azure Static Web Apps and FastAPI runs on Azure Container Apps. Original PDFs are stored privately in Supabase.